

# Mudcard

- **If there is no downside to stratified split but there is a potential downside to not doing it should we just always use it?**
  - Yes, you can always use it.
- **How to sample time-series data which are independent.**
  - I'm not quite sure I understand the question.
  - The level of autocorrelation is an inherent property of time series data, it's not something you or sampling methods have an impact on.
- **I think I'm still trying to understand how the number of K-folds impacts the data. I understand it's important, and I wonder if you increase K to 100 how does this change the model.**
  - It depends on the dataset size.
  - If you have a small dataset, a large k means that you'll have a small number of points in each fold.
  - If you have a large dataset, you'll have enough points in each fold but you'll need to train a large number of models on a large dataset which will be computationally expensive.
- **When is it advantageous to use Kfold splitting versus any 60-20-20 split?**
  - When you use the 60-20-20 split, data is not used efficiently because only 60% of the data is used for training.
  - When you set aside 20% for testing, and use 4 folds on the rest of the data, each fold contains 20% of the data, and you use all 4 folds to train models so you use 80% of the data for training.
  - So kfold splitting is better if you have a small dataset and you want to make most of the small dataset.
  - 60-20-20 split is better for larger datasets because you only train one model, it's more computationally efficient. While when you use k fold, you'll train k models which can be costly if your dataset is large.
- **I am still confused on how the data splitting works in code for the 60:20:20 example for Quiz 1. Why was the train set allocated as 0.75**
  - Work out the sizes of each set and you'll see. :)
  - 25% of 80% is 20%
- **What exactly are classes in the context of splitting a data set (say using k-fold)? I understand that you need to represent a minority class when you split and I see in the lecture that the colors represent different classes, but what are these classes exactly?**
  - In classification, your target variable is categories. For example, in fraud detection, transactions can be real and fraudulent. The class label of a real transaction would be 0, and fraudulent transactions would have a class label of 1.
- **In what instances would an imbalanced dataset be more ideal than a balanced dataset?**
  - It is an inherent property of classification datasets how imbalanced they are. It's not a question of what's ideal for you, it's a question of what the level of imbalance is.
- **Are there specific models that are more suited towards imbalanced dataset?**
  - Yes, we will cover techniques best suited to deal with imbalanced classification problems.

## Split non-iid data

By the end of this lecture, you will be able to

- split non-iid data based on group ID
- split non-iid time series data

## Examples of non-iid data

- if there is any sort of time or group structure in your data, it is likely non-iid
  - **group structure:**
    - **samples are not identically distributed,  $D$  might be different for each group**
    - a person appears multiple times in the dataset (e.g., hospital/doctor visits)
    - data is collected on multiple instruments (e.g., equipment failure prediction)
    - geographical data (e.g., data collected about various cities, counties, states, countries)
  - **time series data**
    - **values are not independent**
    - stocks price
    - covid19 cases
    - weather data

## Ask yourself these questions!

- What is the intended use of the model? What is it supposed to do/predict?
- What data do you have available at the time of prediction?
- Your split must mimic the intended use of the model only then will you accurately estimate how well the model will perform on previously unseen points (generalization error).

- two examples:
  - if you want to predict the outcome of a new patient's visit to the ER:
    - your test score must be based on patients not included in training and validation
    - your validation score must be based on patients not included in training
    - points of one patient should not be distributed over multiple sets because your generalization error will be off
  - if your dataset contains time series features, make sure you don't use future info to predict the past, that's information leakage

## Split non-iid data

By the end of this lecture, you will be able to

- **split non-iid data based on group ID**
- split non-iid time series data

## An example: seizure project

- you can read the publication [here](#)
- classification problem:
  - epileptic seizures vs. non-epileptic psychogenic seizures
- data from empatica wrist sensor
  - heart rate, skin temperature, EDA, blood volume pressure, acceleration
- data collection:
  - patients come to the hospital for a few days
  - eeg and video recording to determine seizure type
  - wrist sensor data is collected
- question:
  - Can we use the wrist sensor data to differentiate the two seizure types on new patients?

```
In [1]: import pandas as pd
import numpy as np

df = pd.read_csv('data/seizure_data.csv')
print(df[df['patient ID'] == 32])
```

	patient ID	seizure_ID	ACC_mean	BVP_mean	EDA_mean	HR_mean	\
5	32	ID32__day3_arm_1_sz1	1.028539	-0.092102	0.112795	64.748167	
6	32	ID32__day3_arm_1_sz1	1.027986	0.745437	0.130486	63.715667	
7	32	ID32__day2_arm_1_sz0	1.002146	0.150810	0.189272	61.838500	
8	32	ID32__day2_arm_1_sz0	1.005410	0.482859	1.226038	66.240833	
9	32	ID32__day1_arm_1_sz0	0.997017	-0.925122	0.200990	56.103667	
10	32	ID32__day1_arm_1_sz0	1.009207	1.618456	1.679754	64.668167	
27	32	ID32__day1_arm_1_sz0	1.000290	0.046690	0.123165	54.289500	
28	32	ID32__day1_arm_1_sz0	1.010351	0.125039	0.471180	65.060667	
29	32	ID32__day2_arm_1_sz0	1.018163	0.254302	0.206010	61.875833	
30	32	ID32__day2_arm_1_sz0	1.016785	1.242893	0.954649	66.216167	
34	32	ID32__day3_arm_1_sz1	1.008867	0.070180	0.195966	65.995667	
35	32	ID32__day3_arm_1_sz1	1.009554	0.222872	0.229909	63.871000	
58	32	ID32__day3_arm_1_sz0	1.008873	-0.550857	0.177822	67.750833	
79	32	ID32__day3_arm_1_sz0	1.026840	0.355953	0.205273	69.124667	

	TEMP_mean	ACC_stdev	BVP_stdev	EDA_stdev	...	BVP_50th	EDA_50th	\
5	36.944833	0.007469	36.486091	0.003905	...	1.815	0.112710	
6	36.676333	0.028190	84.964155	0.018598	...	2.210	0.131921	
7	38.600333	0.003747	64.194294	0.024278	...	6.985	0.186026	
8	39.296083	0.035257	165.665784	0.891139	...	1.140	1.062333	
9	34.656667	0.022648	77.013336	0.132008	...	3.800	0.142159	
10	34.678000	0.046047	146.515297	0.438236	...	5.585	1.690537	
27	38.467417	0.019826	51.176639	0.014530	...	7.765	0.124259	
28	38.448000	0.077142	61.205657	0.156170	...	3.290	0.510114	
29	37.681583	0.006805	40.982246	0.017099	...	1.455	0.202632	
30	37.979500	0.032493	219.277839	0.612229	...	-5.785	1.028171	
34	40.659458	0.021812	49.981175	0.013259	...	3.480	0.198570	
35	40.481333	0.048531	37.409681	0.031963	...	0.695	0.228677	
58	39.906667	0.021431	27.472002	0.003085	...	1.955	0.178073	
79	34.490167	0.008165	40.742936	0.003550	...	3.090	0.206207	

	HR_50th	TEMP_50th	ACC_75th	BVP_75th	EDA_75th	HR_75th	TEMP_75th	\
5	65.060	36.95	1.029947	16.3725	0.115591	65.8175	36.990	
6	62.175	36.81	1.029947	21.1625	0.147611	66.2100	36.840	
7	61.840	38.61	1.006085	43.8850	0.209086	61.9000	38.790	
8	62.325	39.37	1.008872	49.4325	2.313129	71.0625	39.390	
9	56.110	34.66	0.996821	35.2700	0.176739	56.6050	34.660	
10	65.790	34.66	1.021497	70.4800	1.998868	67.7725	34.735	
27	53.960	38.49	1.002073	39.8525	0.133226	54.7425	38.500	
28	65.285	38.45	1.014302	25.4625	0.577047	69.4975	38.530	
29	61.910	37.68	1.022811	29.2125	0.219282	61.9300	37.750	
30	64.700	38.00	1.022811	65.5000	1.503002	69.5725	38.030	
34	66.145	40.68	1.013700	13.1300	0.199852	67.0425	40.710	
35	64.395	40.49	1.016106	12.9650	0.260383	65.9625	40.530	
58	68.170	39.93	1.015264	17.8625	0.179354	68.5725	40.030	
79	69.810	34.37	1.033260	13.4550	0.207488	70.0000	34.680	

	label
5	0.0
6	0.0
7	0.0
8	0.0
9	0.0
10	0.0
27	0.0
28	0.0
29	0.0
30	0.0
34	0.0
35	0.0
58	0.0
79	0.0

[14 rows x 48 columns]

```
In [2]: y = df['label']
patient_ID = df['patient ID']
seizure_ID = df['seizure_ID']
X = df.drop(columns=['patient ID','seizure_ID','label'])
classes, counts = np.unique(y,return_counts=True)
print(classes, counts)
print('balance:',np.max(counts/len(y)))
```

```
[0. 1.] [ 86 190]
balance: 0.6884057971014492
```

```
In [3]: from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import make_scorer

def ML_pipeline_kfold_GridSearchCV(X,y,random_state,n_folds):
```

```

# create a test set
X_other, X_test, y_other, y_test = train_test_split(X, y, test_size=0.2, random_state = random_state, stratify=y)
# splitter for _other
kf = StratifiedKFold(n_splits=n_folds, shuffle=True, random_state=random_state)
# create the pipeline: preprocessor + supervised ML method
scaler = StandardScaler()
pipe = make_pipeline(scaler, SVC())
# the parameter(s) we want to tune
param_grid = {'svc__C': np.logspace(-3, 4, num=8), 'svc__gamma': np.logspace(-3, 4, num=8)}
# prepare gridsearch
grid = GridSearchCV(pipe, param_grid=param_grid, scoring = make_scorer(accuracy_score),
                    cv=kf, return_train_score = True)

# do kfold CV on _other
grid.fit(X_other, y_other)
return grid, grid.score(X_test, y_test)

```

```

In [4]: test_scores = []
for i in range(5):
    grid, test_score = ML_pipeline_kfold_GridSearchCV(X, y, i*42, 5)
    print(grid.best_params_)
    print('best CV score:', grid.best_score_)
    print('test score:', test_score)
    test_scores.append(test_score)
print('test accuracy:', np.around(np.mean(test_scores), 2), '+/-', np.around(np.std(test_scores), 2))

```

```


{'svc__C': 1.0, 'svc__gamma': 0.01}
best CV score: 0.9227272727272726
test score: 0.9285714285714286
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9363636363636363
test score: 0.9285714285714286
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9045454545454547
test score: 0.9464285714285714
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9
test score: 0.9285714285714286
{'svc__C': 10.0, 'svc__gamma': 0.01}
best CV score: 0.9363636363636363
test score: 0.9107142857142857
test accuracy: 0.93 +/- 0.01

```

## This is wrong! A very bad case of data leakage!

- the textbook case of data/information leakage!
- if we just do KFold CV blindly, the points from the same patient end up in different sets
  - when you deploy the model and apply it to data from new patients, that patient's data will be seen for the first time
- the ML pipeline needs to mimic the intended use of the model!
  - we want to split the points based on the patient ID!
  - we want all points from the same patient to be in either train/CV/test

## Group-based split: GroupKFold



```

In [5]: from sklearn.model_selection import GroupKFold
from sklearn.model_selection import GroupShuffleSplit
def ML_pipeline_groups_GridSearchCV(X, y, groups, random_state, n_folds):
    # create a test set based on groups
    splitter = GroupShuffleSplit(n_splits=1, test_size=0.2, random_state=random_state)
    for i_other, i_test in splitter.split(X, y, groups):
        X_other, y_other, groups_other = X.iloc[i_other], y.iloc[i_other], groups.iloc[i_other]
        X_test, y_test, groups_test = X.iloc[i_test], y.iloc[i_test], groups.iloc[i_test]
    # check the split
    # print(pd.unique(groups))
    # print(pd.unique(groups_other))

```

```
# print(pd.unique(groups_test))
# splitter for _other
kf = GroupKFold(n_splits=n_folds)
# create the pipeline: preprocessor + supervised ML method
scaler = StandardScaler()
pipe = make_pipeline(scaler,SVC())
# the parameter(s) we want to tune
param_grid = {'svc__C': np.logspace(-3,4,num=8), 'svc__gamma': np.logspace(-3,4,num=8)}
# prepare gridsearch
grid = GridSearchCV(pipe, param_grid=param_grid,scoring = make_scorer(accuracy_score),
                    cv=kf, return_train_score = True)
# do kfold CV on _other
grid.fit(X_other, y_other, groups=groups_other)
return grid, grid.score(X_test, y_test)
```

```
In [6]: test_scores = []
for i in range(5):
    grid, test_score = ML_pipeline_groups_GridSearchCV(X,y,patient_ID,i*42,5)
    print(grid.best_params_)
    print('best CV score:',grid.best_score_)
    print('test score:',test_score)
    test_scores.append(test_score)
print('test accuracy:',np.around(np.mean(test_scores),2), '+/-' ,np.around(np.std(test_scores),2))
```


```
{'svc__C': 100.0, 'svc__gamma': 0.001}
best CV score: 0.7184206989247313
test score: 0.6666666666666666
{'svc__C': 1.0, 'svc__gamma': 0.1}
best CV score: 0.6416666666666667
test score: 0.059322033898305086
{'svc__C': 10.0, 'svc__gamma': 0.001}
best CV score: 0.5720073891625616
test score: 0.9390243902439024
{'svc__C': 10.0, 'svc__gamma': 0.001}
best CV score: 0.6928409090909091
test score: 0.43243243243243246
{'svc__C': 10000.0, 'svc__gamma': 0.001}
best CV score: 0.5390099715099714
test score: 0.8901098901098901
test accuracy: 0.6 +/- 0.32
```

## The takeaway

- an incorrect cross validation pipeline gives misleading results
  - usually the model appears to be pretty accurate
  - but the performance is poor when the model is deployed
- this can be avoided by a careful cross validation pipeline
  - think about how your model will be used
  - mimic that future use in CV

## Let's take a look at group splitters using toy datasets

### Group-based split: GroupKFold



```
In [7]: from sklearn.model_selection import GroupKFold
import numpy as np

X = np.ones(shape=(8, 2))
y = np.ones(shape=(8, 1))
groups = np.array([1, 1, 2, 2, 2, 3, 3, 3])

group_kfold = GroupKFold(n_splits=3)
```

```
for train_index, test_index in group_kfold.split(X, y, groups):  
    print("TRAIN:", train_index, "TEST:", test_index)
```

TRAIN: [0 1 2 3 4] TEST: [5 6 7]

TRAIN: [0 1 5 6 7] TEST: [2 3 4]

TRAIN: [2 3 4 5 6 7] TEST: [0 1]

In [8]: `help(GroupKFold)`

Help on class GroupKFold in module sklearn.model\_selection.\_split:

```
class GroupKFold(GroupsConsumerMixin, _BaseKFold)
|   GroupKFold(n_splits=5)
|
|   K-fold iterator variant with non-overlapping groups.
|
|   Each group will appear exactly once in the test set across all folds (the
|   number of distinct groups has to be at least equal to the number of folds).
|
|   The folds are approximately balanced in the sense that the number of
|   samples is approximately the same in each test fold.
|
|   Read more in the :ref:`User Guide <group_k_fold>`.
|
|   For visualisation of cross-validation behaviour and
|   comparison between common scikit-learn split methods
|   refer to :ref:`sphx_glr_auto_examples_model_selection_plot_cv_indices.py`
|
|   Parameters
|   -----
|   n_splits : int, default=5
|       Number of folds. Must be at least 2.
|
|       .. versionchanged:: 0.22
|           ``n_splits`` default value changed from 3 to 5.
|
|   Notes
|   -----
|   Groups appear in an arbitrary order throughout the folds.
|
|   Examples
|   -----
|   >>> import numpy as np
|   >>> from sklearn.model_selection import GroupKFold
|   >>> X = np.array([[1, 2], [3, 4], [5, 6], [7, 8], [9, 10], [11, 12]])
|   >>> y = np.array([1, 2, 3, 4, 5, 6])
|   >>> groups = np.array([0, 0, 2, 2, 3, 3])
|   >>> group_kfold = GroupKFold(n_splits=2)
|   >>> group_kfold.get_n_splits(X, y, groups)
|   2
|   >>> print(group_kfold)
|   GroupKFold(n_splits=2)
|   >>> for i, (train_index, test_index) in enumerate(group_kfold.split(X, y, groups)):
|   ...     print(f"Fold {i}:")
|   ...     print(f"  Train: index={train_index}, group={groups[train_index]}")
|   ...     print(f"  Test:  index={test_index}, group={groups[test_index]}")
|   Fold 0:
|     Train: index=[2 3], group=[2 2]
|     Test:  index=[0 1 4 5], group=[0 0 3 3]
|   Fold 1:
|     Train: index=[0 1 4 5], group=[0 0 3 3]
|     Test:  index=[2 3], group=[2 2]
|
|   See Also
|   -----
|   LeaveOneGroupOut : For splitting the data according to explicit
|       domain-specific stratification of the dataset.
|
|   StratifiedKFold : Takes class information into account to avoid building
|       folds with imbalanced class proportions (for binary or multiclass
|       classification tasks).
|
|   Method resolution order:
|       GroupKFold
|       GroupsConsumerMixin
|       _BaseKFold
|       BaseCrossValidator
|       sklearn.utils._metadata_requests._MetadataRequester
|       builtins.object
|
|   Methods defined here:
|
|   __init__(self, n_splits=5)
|       Initialize self. See help(type(self)) for accurate signature.
|
|   set_split_request(self: sklearn.model_selection._split.GroupKFold, *, groups: Union[bool, NoneType, str] = '$UNC
|   HANGED$') -> sklearn.model_selection._split.GroupKFold from sklearn.utils._metadata_requests.RequestMethod.__get__.<
|   locals>
|       Request metadata passed to the ``split`` method.
|
|       Note that this method is only relevant if
|       ``enable_metadata_routing=True`` (see :func:`sklearn.set_config`).
|       Please see :ref:`User Guide <metadata_routing>` on how the routing
|       mechanism works.
|
|       The options for each parameter are:
```

- ``True``: metadata is requested, and passed to ``split`` if provided. The request is ignored if metadata is not provided.
- ``False``: metadata is not requested and the meta-estimator will not pass it to ``split``.
- ``None``: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- ``str``: metadata should be passed to the meta-estimator with this given alias instead of the original name.

The default (``sklearn.utils.metadata\_routing.UNCHANGED``) retains the existing request. This allows you to change the request for some parameters and not others.

.. versionadded:: 1.3

.. note::

This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a :class:`~sklearn.pipeline.Pipeline`. Otherwise it has no effect.

Parameters

groups : str, True, False, or None, default=sklearn.utils.metadata\_routing.UNCHANGED  
Metadata routing for ``groups`` parameter in ``split``.

Returns

self : object  
The updated object.

split(self, X, y=None, groups=None)  
Generate indices to split data into training and test set.

Parameters

X : array-like of shape (n\_samples, n\_features)  
Training data, where ``n\_samples`` is the number of samples and ``n\_features`` is the number of features.

y : array-like of shape (n\_samples,), default=None  
The target variable for supervised learning problems.

groups : array-like of shape (n\_samples,)
 Group labels for the samples used while splitting the dataset into train/test set.

Yields

train : ndarray  
The training set indices for that split.

test : ndarray  
The testing set indices for that split.

-----  
Data and other attributes defined here:

\_\_abstractmethods\_\_ = frozenset()

-----  
Methods inherited from \_BaseKfold:

get\_n\_splits(self, X=None, y=None, groups=None)  
Returns the number of splitting iterations in the cross-validator.

Parameters

X : object  
Always ignored, exists for compatibility.

y : object  
Always ignored, exists for compatibility.

groups : object  
Always ignored, exists for compatibility.

Returns

n\_splits : int  
Returns the number of splitting iterations in the cross-validator.

-----  
Methods inherited from BaseCrossValidator:

\_\_repr\_\_(self)



```

    Return repr(self).

```

---

Methods inherited from sklearn.utils.\_metadata\_requests.\_MetadataRequester:

```

get_metadata_routing(self)
    Get metadata routing of this object.

    Please check :ref:`User Guide <metadata_routing>` on how the routing
    mechanism works.

    Returns
    -----
    routing : MetadataRequest
        A :class:`~sklearn.utils.metadata_routing.MetadataRequest` encapsulating
        routing information.

```

---

Class methods inherited from sklearn.utils.\_metadata\_requests.\_MetadataRequester:

```

__init_subclass__(**kwargs)
    Set the ``set_{method}_request`` methods.

    This uses PEP-487 [1]_ to set the ``set_{method}_request`` methods. It
    looks for the information available in the set default values which are
    set using ``__metadata_request_*`` class attributes, or inferred
    from method signatures.

    The ``__metadata_request_*`` class attributes are used when a method
    does not explicitly accept a metadata through its arguments or if the
    developer would like to specify a request value for those metadata
    which are different from the default ``None``.

    References
    -----
    .. [1] https://www.python.org/dev/peps/pep-0487

```

---

Data descriptors inherited from sklearn.utils.\_metadata\_requests.\_MetadataRequester:


```

__dict__
    dictionary for instance variables

__weakref__
    list of weak references to the object

```

## Group-based split: GroupShuffleSplit



```

In [9]: from sklearn.model_selection import GroupShuffleSplit

gss = GroupShuffleSplit(n_splits=10, train_size=.8, random_state=0)

for train_idx, test_idx in gss.split(X, y, groups):
    print("TRAIN:", train_idx, "TEST:", test_idx)

```

```

TRAIN: [0 1 2 3 4] TEST: [5 6 7]
TRAIN: [0 1 2 3 4] TEST: [5 6 7]
TRAIN: [2 3 4 5 6 7] TEST: [0 1]
TRAIN: [0 1 2 3 4] TEST: [5 6 7]
TRAIN: [0 1 2 3 4] TEST: [5 6 7]
TRAIN: [0 1 5 6 7] TEST: [2 3 4]
TRAIN: [0 1 5 6 7] TEST: [2 3 4]
TRAIN: [2 3 4 5 6 7] TEST: [0 1]
TRAIN: [2 3 4 5 6 7] TEST: [0 1]
TRAIN: [0 1 5 6 7] TEST: [2 3 4]

```

# Quiz 1

Go back to the GroupKFold example above. What happens when you change n\_splits to 4? Why?

Why could we set the n\_splits argument to 5 in GroupShuffleSplit? Check the manual of both methods to find the answer.

Explain your answer in a couple of sentences!


## Split non-iid data

By the end of this lecture, you will be able to

- split non-iid data based on group ID
- **split non-iid time series data**

## Data leakage in time series data is similar!

- do NOT use information in CV which will not be available once your model is deployed
  - don't use future information!



## Time series data

- stock price, crypto price, covid-19 positive case counts, etc
- simple data structure:

time	observation
t_0	y_0
t_1	y_1
t_2	y_2
...	...
t_i	y_i
...	...
t_{n-1}	y_{n-1}
t_n	y_n

- assumption:
  - the difference between two time points (dt) is constant
  - e.g., 1 minute, 5 minutes, 1 hour, or 1 day

## Autocorrelation

- the correlation of the time series data with a delayed copy of itself
- delay on the x axis, correlation coefficient on the y axis
- if delay = 0, the correlation coefficient is 1
- if the delay is short, autocorrelation can be high
- autocorrelation tends to subside for longer delays
- let's check an example

```
In [13]: import pandas as pd
import matplotlib.pyplot as plt
import matplotlib
import numpy as np


df = pd.read_csv('data/daily-min-temperatures.csv')
print(df.shape)
print(df.head())

plt.figure(figsize=(5,3))

plt.plot(df['Temp'])
plt.xticks(np.arange(len(df['Date']))[:365],df['Date'].iloc[:365],rotation=90)
plt.xlabel('date')
plt.ylabel('temperature [C]')
plt.tight_layout()
plt.show()
```

(3650, 2)

	Date	Temp
0	1981-01-01	20.7
1	1981-01-02	17.9
2	1981-01-03	18.8
3	1981-01-04	14.6
4	1981-01-05	15.8



```
In [14]: # let's create an autocorrelation plot

lags = np.arange(3650)
corr_coefs = np.zeros(3650)


for i in np.arange(len(lags)):
    x = df['Temp'].iloc[i:-1].reset_index(drop=True) # recent observations
    y = df['Temp'].iloc[:-i-1].reset_index(drop=True) # lag-shifted observations
    # the shapes must be the same
    if x.shape != y.shape:
        raise ValueError('shape mismatch!')
    # Pearson correlation multiplied by the fraction of time series used
    corr_coefs[i] = x.corr(y,method='pearson')*x.shape[0]/df['Temp'].shape[0]

plt.figure(figsize=(5,3))
plt.plot(lags[2:],corr_coefs[2:])
plt.ylim([-1,1])
plt.xlabel('lag [days]')
plt.ylabel('Pearson correlation coefficient')
plt.title('my figure')
plt.tight_layout()
plt.show()

# a one-liner
plt.figure(figsize=(5,3))

pd.plotting.autocorrelation_plot(df['Temp'])
plt.title("pandas's autocorrelation plot")
plt.tight_layout()
plt.show()
```

```
/Users/azsom/anaconda3/envs/DATA1030/lib/python3.12/site-packages/numpy/lib/function_base.py:2889: RuntimeWarning: D
egrees of freedom <= 0 for slice
  c = cov(x, y, rowvar, dtype=dtype)
/Users/azsom/anaconda3/envs/DATA1030/lib/python3.12/site-packages/numpy/lib/function_base.py:2748: RuntimeWarning: d
ivide by zero encountered in divide
  c *= np.true_divide(1, fact)
/Users/azsom/anaconda3/envs/DATA1030/lib/python3.12/site-packages/numpy/lib/function_base.py:2748: RuntimeWarning: i
nvalid value encountered in multiply
  c *= np.true_divide(1, fact)
```



# Autoregression: create feature matrix using lag features

- goal:
  - predict what y will be dt in the future
- the target variable and lag features:

feature_1	feature_2	...	feature_m-1	feature m	target variable
y_0	y_1	...	y_m-1	y_m	y_m+1
y_1	y_2	...	y_m	y_m+1	y_m+2
...	...	...	...	...	...
y_i-m	y_i-m+1	...	y_i-2	y_i-1	y_i
...	...	...	...	...	...
y_n-m	y_n-m+1	...	y_n-2	y_n-1	y_n

- the features are shifted with respect to the original observation with a dt lag
- this feature matrix should still be split based on time (e.g., older observations in train, most recent observations in test)

```
In [12]: y = df['Temp']
X = pd.concat([df['Temp'].shift(3),df['Temp'].shift(2),df['Temp'].shift(1)],axis=1)
X.columns = ['lag 3 days','lag 2 days','lag 1 day']
print(X.tail(10))
print(y.tail(10))
```

```

lag 3 days lag 2 days lag 1 day
3640      14.7      15.4      13.1
3641      15.4      13.1      13.2
3642      13.1      13.2      13.9
3643      13.2      13.9      10.0
3644      13.9      10.0      12.9
3645      10.0      12.9      14.6
3646      12.9      14.6      14.0
3647      14.6      14.0      13.6
3648      14.0      13.6      13.5
3649      13.6      13.5      15.7
3640      13.2
3641      13.9
3642      10.0
3643      12.9
3644      14.6
3645      14.0
3646      13.6
3647      13.5
3648      15.7
3649      13.0
Name: Temp, dtype: float64
```

## Things to consider

- lag between the target variable and feature m can be more if you want to predict the observation multiple dt's in the future
- you might also have multiple time series to work with (prices of multiple stock, covid cases in multiple countries, etc)
  - all of those need to be shifted by the same lag relative to the target variable
- due to autocorrelation, the features closer in time to the target variable tend to be more predictive
- how many features should you use?
  - treat the number of features as a hyperparameter

## Special scenarios

- what if dt is not consant and/or each time series have its own non-uniform time?
  - for example you try to predict crypto prices based on stock prices
    - stock prices are available once per hour
    - crypto prices are only available when a trade happens (i.e., some tokens are traded rarely)
- interpolate to a uniform time grid
  - try linear and non-linear interpolation techniques to figure out what works best
  - check out [scipy](#) for more info
  - cubic spline interpolation usually works well
- you might have a mix of time series and non-time series features
  - cvs customer purchase history
    - you know what a customer bought and when - time series part
    - you have info on the customer (gender, race, address, etc) - non-time series part

## Mud card

In [ ]: